


System Calls

© Bharati Vidyapeeth's Institute of Computer Applications and Management, New Delhi-63, by Shalini Singh Jaspal

U1.




Library Functions Vs. System Calls

A library function:

- Ordinary function that resides in a library external to the calling program.
- A call to a library function is just like any other function call.
- The arguments are placed in processor registers or onto the stack, and execution is transferred to the start of the function's code, which typically resides in a loaded shared library.
- A library function may have a one-to-one correspondence with a system-call (read)
Or
- The function may format and present data to a system call (printf)
Or
- The function may not require a system call at all.

© Bharati Vidyapeeth's Institute of Computer Applications and Management, New Delhi-63, by Shalini Singh Jaspal

U1.




Library Functions Vs. System Calls

A system call:

- Is implemented in the Linux kernel and performs basic functions that require kernel level privileges
- All activities related to
 - File management
 - User management
 - Process and Memory management
 are handled by the kernel using these system calls.
- Low-level I/O functions such as open and read are examples of system calls on Linux.
- A program running in user mode switches to kernel mode when a system call is made.

© Bharati Vidyapeeth's Institute of Computer Applications and Management, New Delhi-63, by Shalini Singh Jaspal


U1.



System Calls

- When a program makes a system call, the arguments are packaged up and handed to the kernel, the kernel then, takes over execution of the program until the call completes.
- A system call isn't an ordinary function call, and a special procedure is required to transfer control to the kernel.
- The GNU C library (the implementation of the standard C library provided with GNU/Linux systems) wraps Linux system calls with functions so that you can call them easily.
- The set of Linux system calls forms the most basic interface between programs and the Linux kernel. Each call presents a basic operation or capability.
- Note that a library function may invoke one or more other library functions or system calls as part of its implementation.


© Bharati Vidyapeeth's Institute of Computer Applications and Management, New Delhi-63, by Shalini Singh Jaspal
 UI



System Calls

- When an error occurs in a system call:
 - The return value is generally set to -1
 - A global variable errno is set to a positive integer
- This integer
 - Is associated with an error message
 - Is represented by a symbolic constant
- Library functions (like perror) can be used in order to get the information associated with an error.


© Bharati Vidyapeeth's Institute of Computer Applications and Management, New Delhi-63, by Shalini Singh Jaspal
 UI



File Accessibility and Directories

- access
- stat
- chmod
- chown
- chdir
- chroot
- And more...

© Bharati Vidyapeeth's Institute of Computer Applications and Management, New Delhi-63, by Shalini Singh Jaspal
 UI




access

- Determines check user's permissions for a file .
- Can check
 - any combination of read, write, and execute permission,
 - file's existence.
- Syntax:

```
int access(const char *pathname, int mode);
```
- Arguments
 - Path to the file to check
 - Bitwise or of R_OK, W_OK, and X_OK, corresponding to read, write, and execute permission


© Bharati Vidyapeeth's Institute of Computer Applications and Management, New Delhi-63, by Shalini Singh Jaspal
 UI



access

- return value
 - 0 if the process has all the specified permissions.
 - 1 if the file exists but the calling process does not have the specified permissions. The global variable, `errno` is set to `EACCES` (or `EROFS`, if write permission was requested for a file on a read-only file system).
 - If the second argument is `F_OK`, `access` simply checks for the file's existence.
 - ✓ If the file exists, the return value is 0; if not, the return value is -1 and `errno` is set to `ENOENT`
 - ✓ `errno` may instead be set to `EACCES` if a directory in the file path is inaccessible

© Bharati Vidyapeeth's Institute of Computer Applications and Management, New Delhi-63, by Shalini Singh Jaspal
 UI




stat

- Provides information regarding a file contained on its i-node, such as
 - The file size,
 - The last modification time,
 - Permissions,
 - The owner.
- Syntax:

```
int stat (const char * filename, struct stat *buf);
```
- Parameters:
 - The path to the file and
 - A pointer to a variable of type `struct stat`. Set the second parameter to a valid `stat` structure containing the following information:

© Bharati Vidyapeeth's Institute of Computer Applications and Management, New Delhi-63, by Shalini Singh Jaspal
 UI




```

struct stat {
    dev_t    st_dev;    /* device ID, for device files */
    ino_t    st_ino;    /* inode number */
    mode_t    st_mode;  /* type & rights */
    nlink_t   st_nlink; /* number of hard links */
    uid_t     st_uid;   /* user ID of owner */
    gid_t     st_gid;   /* group ID of owner */
    off_t     st_size;  /* total size, in bytes */
    blksize_t st_blksize; /* blocksize for filesystem I/O */
    blkcnt_t  st_blocks; /* number of blocks allocated */
    time_t    st_atime; /* time of last access */
    time_t    st_mtime; /* time of last modification */
    time_t    st_ctime; /* time of last change in inode */
};

```

© Bharati Vidyapeeth's Institute of Computer Applications and Management, New Delhi-63, by Shalini Singh Jaspal
 UI: #



stat


Return:

- 0: If the call to stat is successful and the fields of the structure are filled with information about that file
- -1 on failure

Few useful fields in struct stat:

- st_mode contains the file's access permissions and encodes the type of the file in higher-order bits.
- st_uid and st_gid contain the IDs of the user and group, respectively, to which the file belongs.
- st_size contains the file size, in bytes.
- st_atime contains the time when this file was last accessed (read or written).
- st_mtime contains the time when this file was last modified


© Bharati Vidyapeeth's Institute of Computer Applications and Management, New Delhi-63, by Shalini Singh Jaspal
 UI: #



st_mode Macros

- These macros check the value of the st_mode field value to figure out what kind of file stat was invoked on.
- A macro evaluates to true if the file is of that type.
 - S_ISBLK (mode) block device
 - S_ISCHR (mode) character device
 - S_ISDIR (mode) directory
 - S_ISFIFO (mode) fifo (named pipe)
 - S_ISLNK (mode) symbolic link
 - S_ISREG (mode) regular file
 - S_ISSOCK (mode) socket

© Bharati Vidyapeeth's Institute of Computer Applications and Management, New Delhi-63, by Shalini Singh Jaspal
 UI: #



Flags in *st_mode* for checking Permissions


S_IRWXU 00700 mask for file owner permissions
 S_IRUSR 00400 owner has read permission
 S_IWUSR 00200 owner has write permission
 S_IXUSR 00100 owner has execute permission

S_IRWXG 00070 mask for group permissions
 S_IRGRP 00040 group has read permission
 S_IWGRP 00020 group has write permission
 S_IXGRP 00010 group has execute permission

S_IRWXO 00007 mask for permissions for others (not in group)
 S_IROTH 00004 others have read permission
 S_IWOTH 00002 others have write permission
 S_IXOTH 00001 others have execute permission

© Bharati Vidyapeeth's Institute of Computer Applications and Management, New Delhi-63, by Shalini Singh Jaspal

UI




Symbolic Links: lstat

- If you call stat on a symbolic link, stat follows the link and you can obtain the information about the file that the link points to, not about the symbolic link itself.
- This implies that S_ISLNK will never be true for the result of stat. Use the **lstat** function if you don't want to follow symbolic links; this function obtains information about the link itself rather than the link's target.
- If you call lstat on a file that isn't a symbolic link, it is equivalent to stat. Calling stat on a broken link (a link that points to a nonexistent or inaccessible target) results in an error, while calling lstat on such a link does not.
- If you already have a file open for reading or writing, call **fstat** instead of stat. This takes a file descriptor as its first argument instead of a path.

© Bharati Vidyapeeth's Institute of Computer Applications and Management, New Delhi-63, by Shalini Singh Jaspal

UI




chmod

- Purpose: change mode of file
- The file permission bits of the file named specified by *path* or referenced by the file descriptor *fd* are changed to *mode*.
- The **chmod()** system call verifies that the process owner (user) either owns the file specified by *path* (or *fd*), or is the super-user.
- The **chmod()** system call follows symbolic links to operate on the target of the link rather than the link itself.

© Bharati Vidyapeeth's Institute of Computer Applications and Management, New Delhi-63, by Shalini Singh Jaspal

UI




chmod

- Header: `sys/types.h` , `sys/stat.h`
- Syntax


```
int chmod(const char *path, mode_t mode);
```
- Parameters
 - Path: of the file whose *mode* is to be changed
 - Mode: created from *or'd* permission bit masks defined in `<sys/stat.h>`
- Return Value
 - 0: successful completion; -1: unsuccessful; the global variable *errno* is set to indicate the error.

© Bharati Vidyapeeth's Institute of Computer Applications and Management, New Delhi-63, by Shalini Singh Jaspal
 UI: #



Bit masks for chmod

```

#define S_IRWXU 0000700 /* RWX mask for owner */
#define S_IRUSR 0000400 /* R for owner */
#define S_IWUSR 0000200 /* W for owner */
#define S_IXUSR 0000100 /* X for owner */


#define S_IRWXG 0000070 /* RWX mask for group */
#define S_IRGRP 0000040 /* R for group */
#define S_IWGRP 0000020 /* W for group */
#define S_IXGRP 0000010 /* X for group */

#define S_IRWXO 0000007 /* RWX mask for other */
#define S_IROTH 0000004 /* R for other */
#define S_IWOTH 0000002 /* W for other */
#define S_IXOTH 0000001 /* X for other */

#define S_ISUID 0004000 /* set user id on execution */
#define S_ISGID 0002000 /* set group id on execution */
#define S_ISVTX 0001000 /* save swapped text even after use */

```


© Bharati Vidyapeeth's Institute of Computer Applications and Management, New Delhi-63, by Shalini Singh Jaspal
 UI: #



chown

- Purpose: change owner and group of a file
- The `chown()` system call changes the user and group ownership of a file.
- Only a user having appropriate privileges can change the ownership of a file.
- To get UID corresponding to a user name, use `struct passwd getpwnam (const char *name)`

© Bharati Vidyapeeth's Institute of Computer Applications and Management, New Delhi-63, by Shalini Singh Jaspal
 UI: #




chown

- Header: `unistd.h`
- Syntax:


```
int chown(const char *path, uid_t owner, gid_t group);
```
- Parameters:
 - `path` points to the path name of a file
 - owner ID and group ID of new owner & group
- Return Value:
 - 0: Successful completion.
 - 1: Failure. The owner and group of the file remain unchanged. `errno` is set to indicate the error.


© Bharati Vidyapeeth's Institute of Computer Applications and Management, New Delhi-63, by Shalini Singh Jaspal
 UI: #



chdir

- Purpose: change working directory
- Header: `unistd.h`
- Syntax: `int chdir(const char *path);`
- Parameters: the new directory path
- Return Value:
 - 0: On success
 - 1: on failure, `errno` is set appropriately.


© Bharati Vidyapeeth's Institute of Computer Applications and Management, New Delhi-63, by Shalini Singh Jaspal
 UI: #



chroot

- changes the apparent disk root directory for the current running process and its children
- This directory will be used for pathnames beginning with `/`.
- The root directory is inherited by all children of the current process.
- A program that is re-rooted to another directory cannot access or name files outside that directory


© Bharati Vidyapeeth's Institute of Computer Applications and Management, New Delhi-63, by Shalini Singh Jaspal
 UI: #



chroot

- Header: `unistd.h`
- Syntax: `int chroot(const char *path);`
- Parameters: path of the new root directory
- Return Value:
 - 0: Successful completion.
 - 1: Failure. The owner and group of the file remain unchanged. `errno` is set to indicate the error.


© Bharati Vidyapeeth's Institute of Computer Applications and Management, New Delhi-63, by Shalini Singh Jaspal
UI



Additional

- Try working with these additional File System related System Calls corresponding to simple operations
 - chdir**: Change to a directory
 - mkdir**: Creates a directory
 - rmdir**: Removes an empty directory
 - rename**: Modifies name in a directory entry


© Bharati Vidyapeeth's Institute of Computer Applications and Management, New Delhi-63, by Shalini Singh Jaspal
UI



Process Control

- running Linux Commands from C
- `fork()`
- the `exec` family
- `wait()`
- `exit()`


© Bharati Vidyapeeth's Institute of Computer Applications and Management, New Delhi-63, by Shalini Singh Jaspal
UI



running Linux Commands from C

- Using the `system` function in the standard C library
 - an easy way to execute a shell command from within a program, much as if the command had been typed into a shell.
 - creates a sub-process running the default shell and hands the command to that shell for execution
 - Dependent upon the system shell `/bin/sh`
 - preferable to use the `fork` and `exec` method for creating processes

© Bharati Vidyapeeth's Institute of Computer Applications and Management, New Delhi-63, by Shalini Singh Jaspal
UI



system

Header: `stdlib.h`

Syntax: `int system (const char *command);`


Parameters:

The *command* is executed by issuing `/bin/sh -c command`,

Return

- returns the exit status of the shell command
- If the shell itself cannot be run, `system` returns 127; if another error occurs, `system` returns -1.


© Bharati Vidyapeeth's Institute of Computer Applications and Management, New Delhi-63, by Shalini Singh Jaspal
UI



Using *fork* and *exec*

- Linux provides,
 - one function `fork`, that makes a child process that is an exact copy of its parent process.
 - another set of functions, the `exec` family,
 - causes a particular process to cease being an instance of one program and to instead become an instance of another program.
- To spawn a new process,
 - One first uses `fork` to make a copy of the current process.
 - Then the user uses `exec` to transform one of these processes into an instance of the program you want to spawn.


© Bharati Vidyapeeth's Institute of Computer Applications and Management, New Delhi-63, by Shalini Singh Jaspal
UI



fork()

- When a program calls fork, a duplicate process, called the *child process*, is created.
- The parent process continues executing the program from the point that fork was called.
- The child process, too, executes the same program from the same place.
- Differentiation between parent and child can be done on basis of:
 - PID of the current process
 - Return value of fork.


© Bharati Vidyapeeth's Institute of Computer Applications and Management, New Delhi-63, by Shalini Singh Jaspal
UI



Child / Parent

- The child process is a new process and therefore has a new process ID, distinct from its parent's process ID.
 - One way for a program to distinguish whether it's in the parent process or the child process is to call getpid.
- However, the fork function provides different return values to the parent and child processes
 - one process "goes in" to the fork call, and two processes "come out," with different return values.
 - The return value in the parent process is the process ID of the child.
 - The return value in the child process is zero.
 - This can be used by the program to tell whether it is now running as the parent or the child process.

© Bharati Vidyapeeth's Institute of Computer Applications and Management, New Delhi-63, by Shalini Singh Jaspal
UI




Example

```

int main()
{
    int pid;
    printf("Before fork...\n");
    pid= fork();
    printf("Process Created\n");
}
  
```

© Bharati Vidyapeeth's Institute of Computer Applications and Management, New Delhi-63, by Shalini Singh Jaspal
UI



The value returned by *fork*

The value returned by `fork` can be used in order to check for the parent or the child process:

-1 in case of Error


```

int main()
{
    int pid;
    printf("Before fork:
        \n");
    pid= fork();
    if (pid==0)
        printf("Child \n");
    else
        printf("Parent \n");
}

```

© Bharati Vidyapeeth's Institute of Computer Applications and Management, New Delhi-63, by Shalini Singh Jaspal

U1




the *exec* family

- The `exec` functions replace the program running in a process with another program.
- When a program calls an `exec` function,
 - that process immediately ceases executing that program and
 - begins executing a new program from the beginning,
 - assuming that the `exec` call doesn't encounter an error.
- Within the `exec` family, there are functions that vary slightly in their capabilities and how they are called.

© Bharati Vidyapeeth's Institute of Computer Applications and Management, New Delhi-63, by Shalini Singh Jaspal

U1




the *exec* family

- Functions that contain the letter *p* in their names (`execvp` and `execvp`)
 - accept a program name and
 - search for a program by that name in the current execution path;
 - functions that don't contain the *p* must be given the full path of the program to be executed.
- Functions that contain the letter *v* in their names (`execv`, `execvp`, and `execve`)
 - accept the argument list for the new program as a NULL-terminated array of pointers to strings.

© Bharati Vidyapeeth's Institute of Computer Applications and Management, New Delhi-63, by Shalini Singh Jaspal


U1



The exec Family

- Functions that contain the letter *l* (execl, execlp, and execl)
 - accept the argument list using the C language's varargs mechanism.
- Functions that contain the letter *e* in their names (execve and exece)
 - accept an additional argument, an array of environment variables.
 - The argument should be a NULL-terminated array of pointers to character strings.
 - Each character string should be of the form "VARIABLE=value".
- Because exec replaces the calling program with another one, it never returns unless an error occurs.

© Bharati Vidyapeeth's Institute of Computer Applications and Management, New Delhi-63, by Shalini Singh Jaspal
 UI: #



execl

```


int execl(const char *path, const char *arg0,
    ..., const char *argn, (char *)0);

int exece(const char *path, const char *arg0,
    ..., const char *argn, (char *)0, char
    *const envp[]);

int execlp(const char *file, const char *arg0,
    ..., const char *argn, (char *)0);

int execlpe(const char *file, const char
    *arg0, ..., const char *argn, (char *)0,
    char *const envp[]);
  
```

© Bharati Vidyapeeth's Institute of Computer Applications and Management, New Delhi-63, by Shalini Singh Jaspal
 UI: #



execl Example Usage

```


int ret;
ret= execl ("/bin/ls", "ls", "-l", (char *)0);

int ret;
char *env={"HOME=/usr/home", "LOGNAME=home", (char *)0};
exece ("/bin/ls", "ls", "-l", (char *)0, env);

int ret;
ret = execlp ("ls", "ls", "-l", (char *)0);
  
```

execl can similarly be used for running any other application

© Bharati Vidyapeeth's Institute of Computer Applications and Management, New Delhi-63, by Shalini Singh Jaspal
 UI: #



execv

```

int execv(const char *path, char *const
  argv[]);


int execve(const char *path, char *const
  argv[], char *const envp[]);

int execvp(const char *file, char *const
  argv[]);

int execvpe(const char *file, char *const
  argv[], char *const envp[]);

```

© Bharati Vidyapeeth's Institute of Computer Applications and Management, New Delhi-63, by Shalini Singh Jaspal
 UI



execv Example Usage

```

int ret; char *cmd[] = { "ls", "-l", (char *)0 };
ret = execv ("/bin/ls", cmd);


int ret;
char *cmd[] = { "ls", "-l", (char *)0 };
char *env[] =
  { "HOME=/usr/home", "LOGNAME=home", (char *)0 };
ret = execve ("/bin/ls", cmd, env);

int ret;
char *cmd[] = { "ls", "-l", (char *)0 };
ret = execvp ("ls", cmd);

```

execv can similarly be used for running any other application


© Bharati Vidyapeeth's Institute of Computer Applications and Management, New Delhi-63, by Shalini Singh Jaspal
 UI



Process States

- Child dies before parent =>
 - Child becomes (Z)ombie till it is not removed from process table
- Parent dies before parent =>
 - Child becomes (O)rphan temporarily
 - Inherited by PID 1
- Can use (S)leep (n) to make a process wait for n seconds.

© Bharati Vidyapeeth's Institute of Computer Applications and Management, New Delhi-63, by Shalini Singh Jaspal
 UI




The Parent Process

- After creating a child process, the parent process may:
 - Wait for the termination of the child process
 - ✓To gather child process' exit status
 - Continue execution, without waiting for child.
- The parent process uses the **wait** system call to wait for the termination of the child process.

© Bharati Vidyapeeth's Institute of Computer Applications and Management, New Delhi-63, by Shalini Singh Jaspal

UI: #



wait


Used by a parent process to wait till a child process doesn't terminate.

Behaviour:

- If the process calling wait doesn't have any child processes
 - ✓wait returns -1
- If the calling process has a Zombie child
 - ✓PID of the Zombie is returned to the parent
 - ✓Zombie child is removed from the process table
- If the calling process has a child that hasn't terminated yet
 - ✓Parent process is suspended till the child doesn't terminate
 - ✓When the child terminates, a signal is received by the parent that resumes its execution.
 - ✓A zombie child, if present, is also removed at this point.

© Bharati Vidyapeeth's Institute of Computer Applications and Management, New Delhi-63, by Shalini Singh Jaspal

UI: #




Exit Status of a Child Process

- Apart from returning PID of a child process wait also tells about the termination status of the child process

pid_t wait (int * status);
- The value of *status can be used in order to check the exit status of the child process.

© Bharati Vidyapeeth's Institute of Computer Applications and Management, New Delhi-63, by Shalini Singh Jaspal

UI: #




Threads

- Threads, like processes are mechanism to allow a program to do more than one thing at a time.
- As with processes,
 - threads appear to run concurrently;
 - the Linux kernel schedules them asynchronously,
 - interrupting each thread from time to time to give others a chance to execute.

© Bharati Vidyapeeth's Institute of Computer Applications and Management, New Delhi-63, by Shalini Singh Jaspal

UI




Threads Basics

- Conceptually, a thread exists within a process.
- Threads are a finer-grained unit of execution than processes. When a program is invoked,
 - Linux creates a new process
 - and in that process creates a single thread, which runs the program sequentially.
 - This thread can create additional threads;
 - all these threads run the same program in the same process,
 - but each thread may be executing a different part of the program at any given time.

© Bharati Vidyapeeth's Institute of Computer Applications and Management, New Delhi-63, by Shalini Singh Jaspal

UI




What Processes Share...

- We've seen how a program can fork a child process.
- The child process is initially running its parent's program, with its parent's virtual memory, file descriptors, and so on copied.
- The child process can modify its memory, close file descriptors, and the like without affecting its parent, and vice versa

© Bharati Vidyapeeth's Institute of Computer Applications and Management, New Delhi-63, by Shalini Singh Jaspal


UI



...Threads Don't

- When a program creates another thread, though, **nothing is copied**. The creating and the created thread **share**
 - the same memory space,
 - file descriptors, and
 - other system resources as the original.
- If one thread changes the value of a variable, for instance, the other thread subsequently will see the modified value.
- Similarly, if one thread closes a file descriptor, other threads may not read from or write to that file descriptor.
- Because a process and all its threads can be executing only one program at a time, if any thread inside a process calls one of the exec functions, all the other threads are ended (the new program may, of course, create new threads).


© Bharati Vidyapeeth's Institute of Computer Applications and Management, New Delhi-63, by Shalini Singh Jaspal
 UI



pthread

- GNU/Linux implements the POSIX standard thread API (known as *pthread*).
- All thread functions and data types are declared in the header file `<pthread.h>`.
- The `pthread` functions are not included in the standard C library. Instead, they are in `libpthread`,
- Thus, one needs to add `-lpthread` to the command line while linking a program that creates threads.


© Bharati Vidyapeeth's Institute of Computer Applications and Management, New Delhi-63, by Shalini Singh Jaspal
 UI



Thread IDs and Function

- Each thread in a process is identified by a *thread ID*. When referring to thread IDs in C or C++ programs, use the type `pthread_t`.
- Upon creation, each thread executes a *thread function*.
- This is just an ordinary function and contains the code that the thread should run.
- When the function returns, the thread exits.


© Bharati Vidyapeeth's Institute of Computer Applications and Management, New Delhi-63, by Shalini Singh Jaspal
 UI



Thread Argument

- On GNU/Linux, thread functions take a single parameter, of type `void*`, and have a `void*` return type. The parameter is the *thread argument*:
- GNU/Linux passes the value along to the thread without looking at it.
- Your program can use this parameter to pass data to a new thread.
- Similarly, your program can use the return value to pass data from an exiting thread back to its creator.

© Bharati Vidyapeeth's Institute of Computer Applications and Management, New Delhi-63, by Shalini Singh Jaspal
UI: #




pthread_create

The `pthread_create` function creates a new thread. The parameters are:

- A pointer to a `pthread_t` variable, in which the thread ID of the new thread is stored.
- A pointer to a *thread attribute* object. This object controls details of how the thread interacts with the rest of the program. If you pass `NULL` as the thread attribute, a thread will be created with the default thread attributes. (Thread attributes are beyond the scope of current discussion)
- A pointer to the thread function. This is an ordinary function pointer, of this type: `void* (*) (void*)`
- A thread argument value of type `void*`. Whatever you pass is simply passed as the argument to the thread function when the thread begins executing.


© Bharati Vidyapeeth's Institute of Computer Applications and Management, New Delhi-63, by Shalini Singh Jaspal
UI: #



##

- A call to `pthread_create` returns immediately, and the original thread continues executing the instructions following the call.
- Meanwhile, the new thread begins executing the thread function.
- Linux schedules both threads asynchronously, and your program must not rely on the relative order in which instructions are executed in the two threads.
- The program in the example that follows, creates a thread that prints `x`'s continuously to standard error.
- After calling `pthread_create`, the main thread prints `o`'s continuously to standard error.

© Bharati Vidyapeeth's Institute of Computer Applications and Management, New Delhi-63, by Shalini Singh Jaspal
UI: #



Thread creation-a simple implementation


```

#include <pthread.h>
#include <stdio.h>
/* Prints x's to stderr. The parameter is unused. Does not
return. */
void* print_xs (void* unused)
{
    while (1) { fputc ('x', stderr); }
    return NULL;
}
/* The main program. */
int main ()
{
    pthread_t thread_id;
    /* Create a new thread. The new thread will run print_xs
function. */
    pthread_create (&thread_id, NULL, &print_xs, NULL);
    /* Print o's continuously to stderr. */
    while (1) { fputc ('o', stderr); }
    return 0;
}

```

© Bharati Vidyapeeth's Institute of Computer Applications and Management, New Delhi-63, by Shalini Singh Jaspal

UI




Exiting From a Thread

- Under normal circumstances, a thread exits in one of two ways. One way, as illustrated previously, is by returning from the thread function.
- The return value from the thread function is taken to be the return value of the thread.
- Alternately, a thread can exit explicitly by calling `pthread_exit`.
- This function may be called from within the thread function or from some other function called directly or indirectly by the thread function.
- The argument to `pthread_exit` is the thread's return value.

© Bharati Vidyapeeth's Institute of Computer Applications and Management, New Delhi-63, by Shalini Singh Jaspal

UI



Further reading...

...clone to create threads / processes

And decide what is shared.

© Bharati Vidyapeeth's Institute of Computer Applications and Management, New Delhi-63, by Shalini Singh Jaspal

UI
